

Importing Required Libraries

```
In [78]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing
from sklearn.cluster import KMeans
from pandas.plotting import parallel_coordinates
from sklearn.metrics import silhouette_samples, silhouette_score
from scipy.cluster.hierarchy import linkage, dendrogram
from sklearn.neighbors import NearestNeighbors
from sklearn import cluster
from sklearn import metrics
```

Loading the dataset into a variable

```
In [2]: customer = pd.read_csv('Wholesale customers data.csv')
customer.head()
```

Out[2]:

	Channel	Region	Fresh	Milk	Grocery	Frozen	Detergents_Paper	Delicassen
0	2	3	12669	9656	7561	214	2674	1338
1	2	3	7057	9810	9568	1762	3293	1776
2	2	3	6353	8808	7684	2405	3516	7844
3	1	3	13265	1196	4221	6404	507	1788
4	2	3	22615	5410	7198	3915	1777	5185

The data set refers to clients of a wholesale distributor. It includes the annual spending in monetary units (m.u.) on diverse product categories. Let's explore the wholesale customers data

Data Preprocessing

```
In [3]: customer.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 440 entries, 0 to 439
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   Channel                440 non-null   int64  
1   Region                 440 non-null   int64  
2   Fresh                  440 non-null   int64  
3   Milk                   440 non-null   int64  
4   Grocery                 440 non-null   int64  
5   Frozen                 440 non-null   int64  
6   Detergents_Paper       440 non-null   int64  
7   Delicassen             440 non-null   int64  
dtypes: int64(8)
memory usage: 27.6 KB
```

The data contains 8 attributes with 440 rows, all of integer data type.

```
In [4]: customer.isnull().sum()
```

```
Out[4]: Channel                0
Region                 0
Fresh                  0
Milk                   0
Grocery                 0
Frozen                 0
Detergents_Paper       0
Delicassen             0
dtype: int64
```

There are no null values in our data!

```
In [5]: customer.duplicated().sum()
```

```
Out[5]: 0
```


There are no duplicates in the data!

Let's see the summary statistics for wholesale customer data

```
In [6]: customer.describe()
```

Out[6]:

	Channel	Region	Fresh	Milk	Grocery	Frozen	Deterg
count	440.000000	440.000000	440.000000	440.000000	440.000000	440.000000	
mean	1.322727	2.543182	12000.297727	5796.265909	7951.277273	3071.931818	:
std	0.468052	0.774272	12647.328865	7380.377175	9503.162829	4854.673333	:
min	1.000000	1.000000	3.000000	55.000000	3.000000	25.000000	
25%	1.000000	2.000000	3127.750000	1533.000000	2153.000000	742.250000	
50%	1.000000	3.000000	8504.000000	3627.000000	4755.500000	1526.000000	
75%	2.000000	3.000000	16933.750000	7190.250000	10655.750000	3554.250000	:
max	2.000000	3.000000	112151.000000	73498.000000	92780.000000	60869.000000	40



We can observe that attributes have high variability


Let's remove the channel and region columns from our data for better analysis of numerical values only.

```
In [7]: customer = customer.drop('Channel', axis=1)
customer = customer.drop('Region', axis=1)
```

Let's visualize our data using box plot and histogram to learn the distribution of data

```
In [8]: sns.boxplot(data = customer, orient = 'h')
```

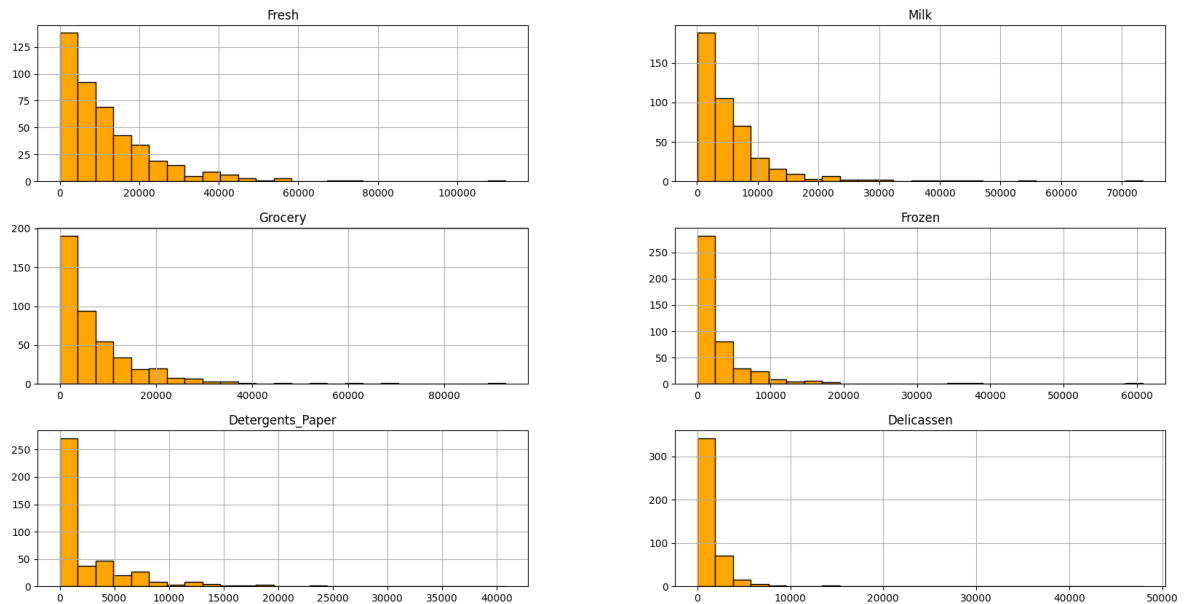
Out[8]: <Axes: >



Let's check the distribution of attributes

```
In [9]: customer.hist(figsize = (20,10), bins = 25, color='orange', edgecolor='black')
```

```
Out[9]: array([[<Axes: title={'center': 'Fresh'}>,  
               <Axes: title={'center': 'Milk'}>],  
              [<Axes: title={'center': 'Grocery'}>,  
               <Axes: title={'center': 'Frozen'}>],  
              [<Axes: title={'center': 'Detergents_Paper'}>,  
               <Axes: title={'center': 'Delicassen'}>]], dtype=object)
```




From the histogram we can observe that the data is not normally distributed. Thus, we will normalize the wholesale customers data.

We will use sci-kit learn library to normalize the data.

```
In [10]: cust_norm = preprocessing.normalize(customer)
cust_norm = pd.DataFrame(cust_norm)
cust_norm.rename(columns = {0:'Fresh',1:'Milk',2:'Grocery',3:'Frozen',4:'Deter',
cust_norm.info()
cust_norm.hist(figsize = (20,10), bins = 25, color='orange', edgecolor='black')
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 440 entries, 0 to 439
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0    Fresh      440 non-null    float64
1    Milk       440 non-null    float64
2    Grocery    440 non-null    float64
3    Frozen     440 non-null    float64
4    Detergents 440 non-null    float64
5    Deli       440 non-null    float64
dtypes: float64(6)
memory usage: 20.8 KB
```

```
Out[10]: array([[<Axes: title={'center': 'Fresh'}>,
<Axes: title={'center': 'Milk'}>],
[<Axes: title={'center': 'Grocery'}>,
<Axes: title={'center': 'Frozen'}>],
[<Axes: title={'center': 'Detergents'}>,
<Axes: title={'center': 'Deli'}>]], dtype=object)
```



This doesn't look like a normal distribution but its an improvement from what we had before. We have completed the preprocessing steps. Now let's apply the K-means clustering algorithm


K-means Clustering

We will apply kmeans clustering from sci-kit learn library. We will apply the number of clusters ranging from 1 to 10 and then select the most appropriate number of clusters using the elbow method

```
In [11]: SSE = []
clusters = [1,2,3,4,5,6,7,8,9,10]
for k in clusters:
    kmeans = KMeans(n_clusters = k, n_init='auto')
    kmeans.fit(cust_norm)
    SSE.append(kmeans.inertia_)

plt.plot(clusters, SSE)
plt.xlabel('Number of Clusters')
plt.ylabel('SSE')
plt.title('Determining the number of clusters')
```

```
Out[11]: Text(0.5, 1.0, 'Determining the number of clusters')
```



Looking at the elbow plot, we can see that 2 number of clusters seems to be the ideal choice for this dataset. So let's fit the algorithm again for 2 clusters.

```
In [12]: kmeans1 = KMeans(n_clusters=2, n_init='auto')
customer_kmeans = kmeans1.fit_predict(cust_norm)
```

Now, let's obtain the labels from our model and check if every data point has been assigned a label or not

```
In [13]: len(cust_norm) == len(kmeans1.labels_)
```

```
Out[13]: True
```

We have labels for every record. Now let's concat labels field to our data and rename the columns

```
In [14]: labels = pd.DataFrame(kmeans1.labels_)
labels.rename(columns = {0:'label'}, inplace = True)
data = pd.concat([cust_norm, labels], axis=1)
```

```
In [15]: data.head()
```

```
Out[15]:
```

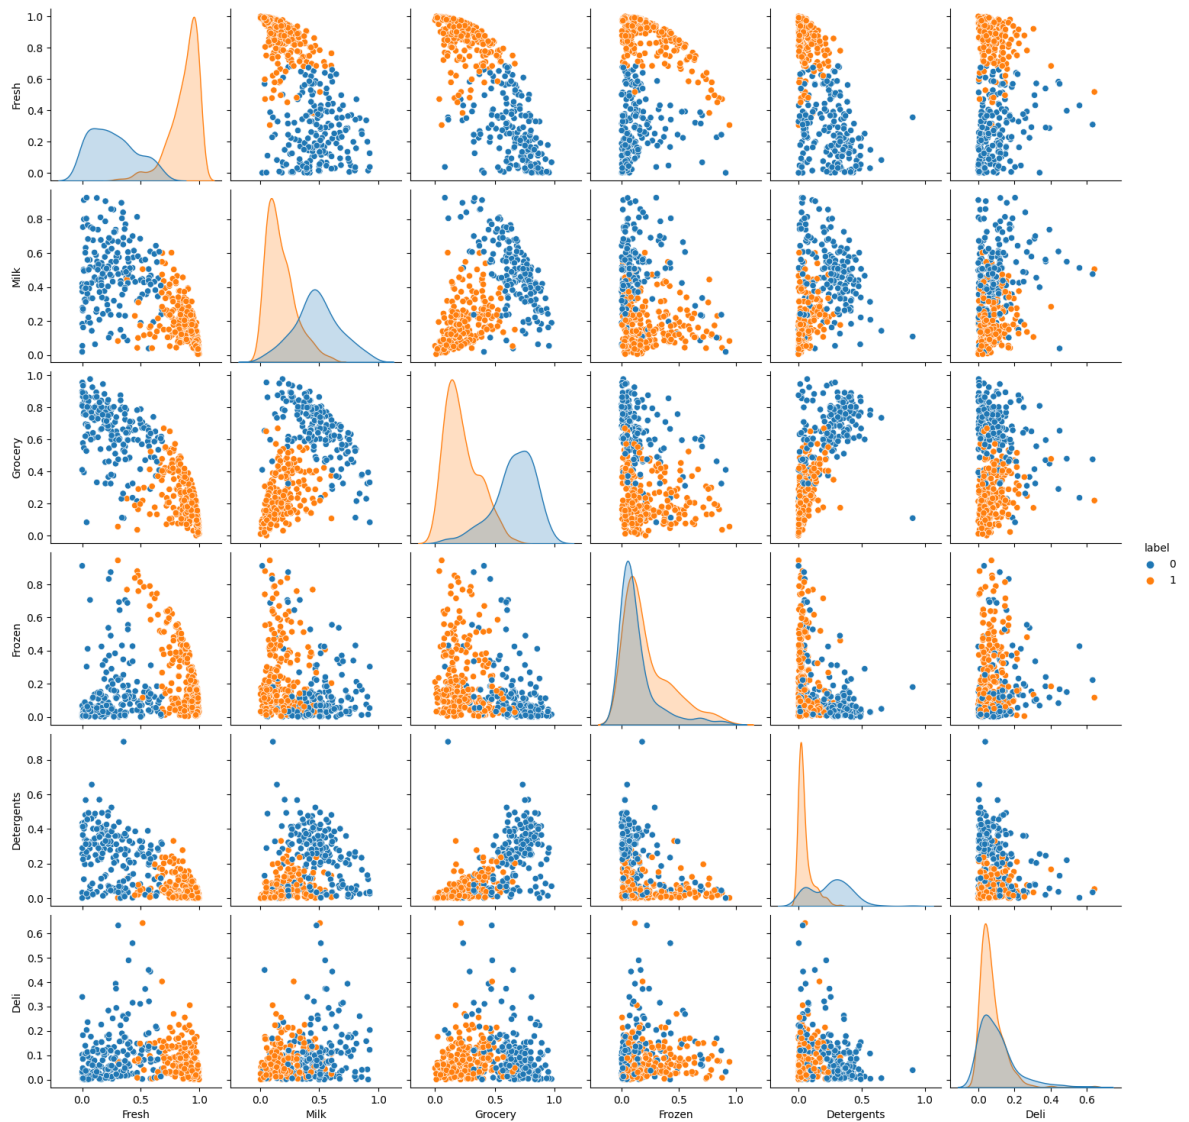
	Fresh	Milk	Grocery	Frozen	Detergents	Deli	label
0	0.708333	0.539874	0.422741	0.011965	0.149505	0.074809	1
1	0.442198	0.614704	0.599540	0.110409	0.206342	0.111286	0
2	0.396552	0.549792	0.479632	0.150119	0.219467	0.489619	0
3	0.856837	0.077254	0.272650	0.413659	0.032749	0.115494	1
4	0.895416	0.214203	0.284997	0.155010	0.070358	0.205294	1

Visualizing Clusters

Finally, let's visualize our clustering model using scatterplots



```
In [16]: sns.pairplot(data, hue = 'label')
```

```
Out[16]: <seaborn.axisgrid.PairGrid at 0x1e326451950>
```



This is bit difficult to interpret, let's plot a parallel coordinates plot for the clusters to obtain better understanding of results.


```
In [17]: plt.figure(figsize=(8,4))
parallel_coordinates(data, 'label', color=('coral','darkblue'))
plt.title('Parallel Coordinates plot for clusters')
plt.ylabel('Annual Spending')
plt.xlabel('Categories')
plt.show()
```



We will also plot the parallel coordinates plot for the centroids of the clustering result.

```
In [18]: # Create a dataframe containing centroids value from clustering output
centroids = pd.DataFrame(kmeans1.cluster_centers_, columns = data.columns[0:6])
centroids['cluster'] = centroids.index
```

```
In [109]: plt.figure(figsize=(8,4))
parallel_coordinates(centroids, 'cluster', color=('darkblue','coral'))
plt.title('Parallel Coordinates plot for cluster centroids')
plt.ylabel('Annual Spending')
plt.xlabel('Categories')
plt.show()
```



Silhouette Analysis

```
In [20]: n_clusters = [2,3,4,5,6,7,8,9,10]
for k in n_clusters:
    sil_kmeans = KMeans(n_clusters = k, n_init='auto')
    cluster_labels = sil_kmeans.fit_predict(cust_norm)

    #The silhouette score gives the average value for all the samples
    sil_avg = silhouette_score(cust_norm,cluster_labels)
    print("For n_cluster =",k,"The average silhouette score is: ",sil_avg)
```

```
For n_cluster = 2 The average silhouette score is: 0.5002248259665941
For n_cluster = 3 The average silhouette score is: 0.4365632328906848
For n_cluster = 4 The average silhouette score is: 0.3798290278483874
For n_cluster = 5 The average silhouette score is: 0.3747453809922325
For n_cluster = 6 The average silhouette score is: 0.3643696320545916
For n_cluster = 7 The average silhouette score is: 0.3195688089575444
For n_cluster = 8 The average silhouette score is: 0.3329818912816675
For n_cluster = 9 The average silhouette score is: 0.2680799795960205
For n_cluster = 10 The average silhouette score is: 0.3141539002326934
```


Thus, the highest silhouette score is obtained for $n_clusters = 2$.

DBSCAN

Let's pick min_samples = 20 and We will determine the eps value using knn

```
In [110]: neighbors = NearestNeighbors(n_neighbors=20)
neighbors_fit = neighbors.fit(customer)
distances, indices = neighbors_fit.kneighbors(customer) #will find k neighbors
```

```
In [111]: distances = np.sort(distances, axis=0)
distances = distances[:,1]
plt.plot(distances)
plt.ylim(0,9000)
plt.xlim(100,600)
plt.show()
```




We determine for minimum 20 samples in a cluster the eps value to be 5000

```
In [112]: dbscan = cluster.DBSCAN(eps=6000, min_samples=2)
clustering_labels = dbscan.fit_predict(customer.to_numpy())
```

```
In [113]: customer['db_label'] = clustering_labels
```

```
In [114]: plt.figure(figsize=(8,4))
parallel_coordinates(customer, 'db_label', color=('coral','darkblue'))
plt.title('Parallel Coordinates plot for clusters')
plt.ylabel('Annual Spending')
plt.xlabel('Categories')
plt.show()
```



```
In [115]: metrics.silhouette_score(customer, customer['db_label'])
```

Out[115]: 0.14647081024822362

In []: